

# CS 564 Project Stage 5

## Front-End and Database Utilities

You will not do this stage. Your instructor will give the precise instruction.

### Introduction

For this part of the project, you will implement the front-end and database utilities of Minirel. You will also have to design and implement the Minirel catalog relations. To help you implement the front-end quickly, we will provide a parser to parse user commands and (a subset of ) SQL. The parser works by reading in one line from stdin, parsing the line and making the appropriate calls to the back-end.

### Database Utilities

There are three database utilities each of which can be run as a standalone UNIX program

- `dbcreate <dbname>`
- `dbdestroy <dbname>`
- `minirel <dbname>`

We have provided complete implementations of `dbcreate <dbname>` and `dbdestroy <dbname>` and a skeleton of `minirel <dbname>` functions for these three routines in the files `minirel.C`, `dbcreate.C`, and `dbdestroy.C` respectively. Each of these executables can be created by giving appropriate arguments to `make`.

#### 1. `dbcreate <dbname>`

This creates the database `dbname`. You should examine this code carefully as it will give you some hints about how to implement operations like `create relation()`. The program first creates a unix directory to hold the database and then changes the working directory to that directory. Next it creates two relations called `RelCatalog` and `AttrCatalog` for the database catalogs. First of course it creates these two relations by creating instances of the `RelCatalog` and `AttrCatalog` classes. More details about this will be given in the section on Catalogs.

#### 2. `dbdestroy <dbname>`

This destroys the database `dbname`. This is done by using the following UNIX command:

```
ostringstream sysString;  
sysString << "rm -r " << dbname;  
system(sysString.str());
```

By doing this, the directory `dbname` and all the files (relations) it contains will be deleted.

#### 3. `minirel <dbname>`

This is the main routine which is used to access and query the Minirel database *dbname*. When *minirel* is invoked with a *dbname*, the very first thing it does is to change its working UNIX directory (using `chdir`) to *dbname*. Each Minirel database corresponds to a UNIX directory and each relation within that database corresponds to a file within the directory. The *minirel* utility starts off by creating a buffer manager, opening the relation and attribute catalogs and then calling the function *parse()*.

**Important: at this point global variable *relCat* refers to the relation catalog table and global variable *attrCat* refers to the attribute catalog table. Keep in mind these two variables, because when you implement various methods of class *RelCatalog*, you may have to refer to the attribute catalog table, and vice versa, when you implement various methods of class *AttrCatalog*, you may have to refer to the relation catalog table.**

With the call to *parse()*, an infinite loop will start prompting the user for a new command, calling the appropriate routines of the lower layers to execute the command and then prompting the user again. The parser (provided by us) will be responsible for generating the prompt. After a command is parsed, it is transformed by the *interp()* function (in file *parser/interp.C*) into a sequence of operators (e.g. selects, joins, creates, destroys, prints, loads etc.), which are then executed one by one by the database back-end. Executing an operator involves calling some routine which you have or will implement in other parts of the project. For example, when the front-end receives a select-project query, the parser will call a procedure named *select()* that will in turn call the methods of the *HeapFileScan* class to scan the appropriate heapfile.

## SQL DDL commands

Once *minirel* has been invoked, it can be used to either answer queries (using SQL DML) or to perform various utility operations (via SQL DDL). (Note: DML = Data Manipulation Language, and DDL = Data Definition Language.) Part5 of the project is mainly concerned with implementing the front-end utilities necessary to support a part of SQL DDL. The syntax of these commands will be described using a pseudo context-free grammar. All commands to the Minirel front-end must end with a semi-colon.

### Front-End Command Syntax

```
<SQL_COMMAND> ::= <DDL_STATEMENT>;  
                | <DML_STATEMENT>;
```

```
<DDL_STATEMENT> ::= <CREATE>  
                   | <DESTROY>  
                   | <LOAD>  
                   | <PRINT>  
                   | <HELP>  
                   | <QUIT>
```

We now describe the semantics of each utility in detail.

**<CREATE> ::= CREATE TABLE *relName* (*attr1 format1*, ..., *attrN formatN*) ;**

The *create* command creates the relation *relName* with the given schema. Relation and attribute names can be at most MAXNAME = 32 characters long. Each format argument has a type specification ('int', 'real', 'char' or 'char' followed by an integer inside a pair of parenthesis specifying the length of the string. e.g. 'char(20)'). For the 'char' format valid lengths are between 1 to 255 (both inclusive). This command results in an invocation of the *RelCatalog::createRel()* method that updates the catalogs appropriately and then creates a *HeapFile* to store the tuples of the relation. *Thus, once you have implemented the method RelCatalog::createRel(), parse() will call that method to execute this command.*

**<DESTROY> ::= DESTROY TABLE *relName* ;**

The *destroy* command should destroy the HeapFile for relation relName. Finally it should remove all relevant catalog entries for that relation. All this is done by invoking the RelCatalog::destroyRel method. *Thus, once you have implemented the method RelCatalog::destroyRel, parse() will call that method to execute this command.*

**<PRINT> ::= PRINT TABLE relName**

The *print* routine prints the contents of the specified relation in a reasonably tidy format. This is done by the UT\_Print function. *Thus, once you have implemented the function UT\_Print, parse() will call that function to execute this command.*

**<LOAD> ::= LOAD TABLE relname FROM (filename) ;**

The load utility bulk loads tuples into a relation (which must have already been created) using data in a binary UNIX file named filename. The load command begins by examining the relation catalog to determine length of the tuple length. Then it opens the UNIX file filename containing the input data. Each input tuple is in binary format (i.e. there is no space wasted for delimiters). The next step is to create an instance of the InsertFileScan class, passing the relname as parameter which cause the appropriate HeapFile to be opened. Finally, one tuple at a time should be read from the UNIX file and inserted into the HeapFile using the InsertFileScan::insertRecord() method. After loading all the tuples, the UNIX file should be closed and the InsertFileScan object should be deleted. This will result in the underlying DB files being closed. All this happens in the UT\_Load function (described below). *Thus, once you have implemented the function UT\_Load, parse() will call that function to execute this command.*

**<HELP> ::= HELP [TABLE relName] ;**

If a relname is not specified, help lists the names of all the relations in the database. Otherwise, it lists the name, type, length and offset of each attribute together with any other information you feel is useful. This is done by the RelCatalog::help function. *Thus, once you have implemented the function RelCatalog::help, parse() will call that function to execute this command.*

**<QUIT> ::= QUIT ;**

The quit command exits minirel, causing all buffers and relations to be flushed to disk. Prior to terminating, the system should close down any open files (including catalogs). This is achieved by the UT\_Quit function. *Thus, once you have implemented the function UT\_Quit, parse() will call that function to execute this command.*

**Summary: To implement the above commands, you will have to implement RelCatalog::createRel(), RelCatalog::destroyRel, RelCatalog::help, UT\_Print, UT\_Load, and UT\_Quit, among others. These methods and functions are described below.**

## Implementing Catalogs

One of the neat things about a relational database is that in addition to the data, even the metadata that describe what relations exist in the database, the types of their attributes and so on are also stored in relations, called catalogs. Minirel has two catalog relations : *relcat* and *attrcat*. The relcat relation contains one tuple for every relation in the database (including itself). The attrcat relation contains one tuple for every attribute of every relation (including the catalog relations) and this tuple contains information about the attribute. Both attrcat and relcat are created by the dbcreate utility and together they contain the schema of the database. Since each database has its own schema, relcat and attrcat must be placed in the UNIX directory corresponding to the database. relcat and attrcat are instances of the RelCatalog and AttrCatalog classes respectively.

## The RelCatalog Class

```

#define RELCATNAME "relcat"      // name of relation catalog
#define ATTRCATNAME "attrcat"   // name of attribute catalog
#define MAXNAME 32              // length of relName, attrName
#define MAXSTRINGLEN 255        // max length of string attribute

// schema for tuples in the relation relcat:
// relation name : char(32) <-- lookup key
// attribute count : integer(4)
typedef struct {
    char relName[MAXNAME];      // relation name
    int attrCnt;                // number of attributes
} RelDesc;

class RelCatalog : public HeapFile {
public:
    // open relation catalog
    RelCatalog(Status &status);

    // get relation descriptor for a relation
    const Status getInfo(const string & relName, RelDesc& record);

    // add information to catalog
    const Status addInfo(RelDesc & record);

    // remove tuple from catalog
    const Status removeInfo(const string & relName);

    // create a new relation
    const Status createRel(
        const string & relName,
        const int attrCnt, // number of elements in attrList[]
        const attrInfo attrList[] // see defn of attrInfo below
    );

    // define attrInfo
    typedef struct {
        char relName[MAXNAME]; // relation name
        char attrName[MAXNAME]; // attribute name
        int attrType; // INTEGER, FLOAT, or STRING
        int attrLen; // length of attribute in bytes
        void *attrValue; // ptr to binary value
    } attrInfo;

    // destroy a relation
    const Status destroyRel(const string & relName);

    // print catalog information (relation may be an empty string)
    const Status help(const string & relName);

    // get rid of catalog
    ~RelCatalog();
};

```

#### **RelCatalog::RelCatalog()**

The constructor just invokes the HeapFile constructor with the constant RELCATNAME (declared above).

## **RelCatalog::~RelCatalog()**

You need to do nothing here as the underlying HeapFile will be close automatically.

## **const Status createRel(const string & relName, const int attrCnt, const attrInfo attrList[])**

First make sure that a relation with the same name doesn't already exist (by using the getInfo() function described below). Next add a tuple to the relcat relation. Do this by filling in an instance of the RelDesc structure above and then invoking the RelCatalog::addInfo() method. Third, for each of the attrCnt attributes, invoke the AttrCatalog::addInfo() method of the attribute catalog table (**remember that this table is referred to by the global variable attrCat**), passing the appropriate attribute information from the attrList[] array as an instance of the AttrDesc structure (see below). Finally, create a HeapFile instance to hold tuples of the relation (hint: there is a procedure to do this which we have seen in the last project stage; you need to give it a string that is the relation name). *Implement this function in create.C*

## **const Status destroyRel(const string & relName)**

First remove all relevant schema information from both the relcat and attrcat relations. Then destroy the heapfile corresponding to the relation (hint: again, there is a procedure to destroy heap file that we have seen in the last project stage; you need to give it a string that is the relation name). *Implement this function in destroy.C*

## **const Status addInfo(RelDesc & record)**

Adds the relation descriptor contained in *record* to the relcat relation. RelDesc represents both the in-memory format and on-disk format of a tuple in relcat. First, create an InsertFileScan object on the relation catalog table. Next, create a record and then insert it into the relation catalog table using the method insertRecord of InsertFileScan.

## **const Status getInfo(const string & relName, RelDesc& record)**

Open a scan on the relcat relation by invoking the startScan() method on itself. You want to look for the tuple whose first attribute matches the string relName. Then call scanNext() and getRecord() to get the desired tuple. Finally, you need to memcpy() the tuple out of the buffer pool into the return parameter *record*.

## **const Status removeInfo(const string & relName)**

Remove the tuple corresponding to relName from relcat. Once again, you have to start a filter scan on relcat to locate the rid of the desired tuple. Then you can call deleteRecord() to remove it.

## **const Status help(const string & relName)**

If relation.empty() is true (empty() is a method on the string class), print (to standard output) a list of all the relations in relcat (including how many attributes it has). Otherwise, print all the tuples in attrcat that are relevant to relName. *Implement this function in help.C*

## **The AttrCatalog Class**

```
// schema of tuples in the attribute catalog:
typedef struct {
char relName[MAXNAME]; // relation name
char attrName[MAXNAME]; // attribute name
int attrOffset; // attribute offset
int attrType; // attribute type
int attrLen; // attribute length
} AttrDesc;
```

```

class AttrCatalog : public HeapFile {
friend RelCatalog;

public:
// open attribute catalog
AttrCatalog(Status &status);

// get attribute catalog tuple
const Status getInfo(const string & relName,
                    const string & attrName,
                    AttrDesc &record);

// add information to catalog
const Status addInfo(AttrDesc & record);

// remove tuple from catalog
const Status removeInfo(const string & relName,
                      const string & attrName);

// get all attributes of a relation
const Status getRelInfo(const string & relName,
                      int &attrCnt,
                      AttrDesc *&attrs);

// delete all information about a relation
const Status dropRelation(const string & relName);

// close attribute catalog
~AttrCatalog();
};

```

#### **AttrCatalog(Status &status)**

Just call the new HeapFile constructor (see below), passing the constant ATTRCATNAME (defined above)

#### **~AttrCatalog()**

Do nothing. File gets closed automatically by HeapFile destructor.

#### **const Status getInfo(const string & relName, const string & attrName, AttrDesc &record)**

Returns the attribute descriptor record for attribute attrName in relation relName. Uses a scan over the underlying heapfile to get all tuples for relation and check each tuple to find whether it corresponds to attrName. (Or maybe do it the other way around !) This has to be done this way because a predicated HeapFileScan does not allow conjuncted predicates. Note that the tuples in attrcat are of type AttrDesc (structure given above).

#### **const Status getRelInfo(const string & relName, int &attrCnt, AttrDesc \*&attrs)**

While getInfo() above returns the description of a single attribute, this method returns (by reference) descriptors for all attributes of the relation via *attr*, an array of AttrDesc structures, and the count of the number of attributes in *attrCnt*. The attrs array is allocated by this function, but it should be deallocated by the caller.

#### **const Status addInfo(AttrDesc & record)**

Adds a tuple (corresponding to an attribute of a relation) to the attrcat relation.

**const Status removeInfo(const string & relName, const string & attrName)**

Removes the tuple from attrcat that corresponds to attribute attrName of relation.

**const Status dropRelation(const string & relName)**

Deletes all tuples in attrcat for the relation relName. Again use a scan to find all tuples whose relName attribute equals relation and then call deleteRecord() on each one. *Implement this function in destroy.C*

## Implementing Utility Routines

When the Minirel parser reads in a command, it will parse the command and call the necessary function to execute it. Some of the front-end DDL commands such as create and destroy are implemented by invoking the appropriate methods on the catalog classes. The implementation of the other utility functions load, print and quit are described below.

**void UT\_quit(void)**

This function should simply delete the RelCatalog and AttrCatalog objects (causing the underlying files to be closed), delete the BufMgr object (causing any remaining dirty pages to be flushed to disk) and then terminate the minirel program by calling exit().

**const Status UT\_load(const string & relName, const string & fileName)**

Loads the tuples from the UNIX file fileName into the specified relation. The relation must already exist before UT\_load is invoked.

**const Status UT\_print(const string & relName)**

Prints the values of all the attributes of all tuples of the specified relation. We will give you this code.

## Error Handling

You should continue to use the same error handling mechanism that you used for the previous parts of the project. Feel free to augment this with new error codes and messages as needed.

## Getting Started

Start by downloading the Zip file for this stage (**since you are not doing this stage, you can just ignore this**). This Zip file contains the following files that are relevant to this part of the project (in addition to other files which were created while developing the lower layers).

- minirel.C - Main program for Minirel. Can be used without modification.
- dbcreate.C, dbdestroy.C - Main programs for these utilities. We have provided complete implementations of these functions.
- create.C, destroy.C - Stubs of catalog functions for creating and dropping relations. **You have to fill these out.**
- catalog.C - Stubs for other catalog functions. **You have to implement these.**
- load.C - Utility for loading a relation from a UNIX file. **You have to implement this.**
- print.C - Utility for printing the data in a relation.
- help.C - Utility for printing out schema. **You have to implement this.**
- quit.C - Utility for cleaning up and exiting Minirel.
- select.C, join.C, insert.C, delete.C - Stubs for functions which will be used later in the project. You should not modify these.

- Other .h files - These contain the relevant class definitions and function prototypes. You need not modify these.
- Makefile - To compile this part of the project.
- Files in the parser subdirectory - This contains the SQL parser and interpreter. You should not modify these.

In the Zip file you will find a shell program called **uttest** along with two directories **uttestqueries** and **uttestdata**. If you look at the files in the directory **uttestqueries** you will notice that they contain a progressively more difficult series of tests. The directory **uttestdata** contains some small tables that will get loaded by the test queries.

To run all the tests, simply invoke **uttest** while in the part 5 directory. If you want to run one test at a time simply type **uttest i** where **i** is a number between 1 and 10. For example, **uttest 3** will run the test in the file **ut.3** in the **uttestqueries** directory.